

## Lecture 6: Processes

**Dr. Naureen Nizam**  
[nnizam@cs.toronto.edu](mailto:nnizam@cs.toronto.edu)

CSCB09H3: Software Tools and Systems Programming (SY 110)  
 Department of Computer and Mathematical Science  
 Feb 11/13, 2019



Tomorrow is  
created here.

UNIVERSITY OF TORONTO SCARBOROUGH  
 1265 Military Trail, Toronto, Ontario M1C 1A4

## Administrivia

- **Assignments**
  - A2 Posted
  - Due: March 3, 2019
- **Labs**
  - Lab 4 solutions posted
  - Lab 5 posted – Memory Leaks & Makefile
- **Reading Week – Next Week (No Class and No Labs)**
- **Midterm**
  - Monday, February 25<sup>th</sup> (7 PM SW143, SW128, SW319)
  - Wednesday, February 27<sup>th</sup> (12 PM SW143) – Makeup (only those who are approved)

2

## Last Week

Helpful  
Tips

- C Programming:

- Tricks and Tips

- Passing arguments to functions
- Passing command line args to program
- typedef
- Organizing your program
- Makefiles
- Pre-processor directives
- Error handling

KING Chapters 13, 14, 15, 16, 17, 22, 24

- File I/O

3

## I/O in C

- Have already seen library functions (from stdio.h) in sample programs:
 

```
int x = 10;
printf("The value of x is %d\n", x);
```
- Printing a character:
 

```
-int putchar(int c);
```
- Printing a string
 

```
-int puts(const char *s);
```
- These all print to stdout (standard output = screen)
  - What if we want to print to a file?

4

## File I/O in C

- Have already seen library functions (from stdio.h) in sample programs:
 

```
int x = 10;
printf("The value of x is %d\n", x);
-int fprintf(FILE *stream, const char *format, ...);
```
- Printing a character:
 

```
-int putchar(int c);
-int fputc(int c, FILE *stream);
```
- Printing a string
 

```
-int puts(const char *s);
-int fputs(const char *s, FILE *stream);
```
- What is FILE \*stream?

5

## File interfaces in C/Unix

- Two main mechanisms for managing file access:
- File descriptors (low-level):
  - Each open file is identified by a small integer
  - Use for pipes, sockets (will see later what those are ...)
- File pointers (regular files):
  - You use a pointer to a file structure (FILE \*) as handle to a file.
  - The file struct contains a file descriptor and a buffer.
  - Use for regular files

6

## Standard streams

- All programs automatically have three files open:

- `FILE *stdin;`
- `FILE *stdout;`
- `FILE *stderr;`

	stdio name	File descriptor
Standard input	<code>stdin</code>	0
Standard output	<code>stdout</code>	1
Standard error	<code>stderr</code>	2

} Comes by default from the keyboard.

} Go by default to the screen.

- Those are ready to use:

```
fprintf(stdout, "Hello!\n"); // Identical to printf("Hello!\n");
fputs("Error!\n", stderr);
```

7

## Operating on regular files

- A file first needs to be opened to obtain a `FILE *`

```
FILE *fopen(const char *filename,
            const char *mode);
```

- `filename` identifies the file to open.
- `mode` tells how to open the file:
  - "r" for reading, "w" for writing, "a" for appending
- returns a pointer to a `FILE` struct which is the handle to the file. This pointer will be used in subsequent operations.
- To close a file: `void fclose(FILE *stream);`

```
FILE *fp = fopen("my_file.txt", "w");
fprintf(fp, "Hello!\n");
fclose(fp);
```

8

## Reading from files

- `char *fgets(char *s, int size, FILE *stream);`

- Reads until a `\n` or `size-1` characters, whichever is first
- Stores characters in buffer `s` points to
- If a newline is read, it is stored into the buffer too.
- Always null-terminates the string

- How would you use `fgets` to read from the keyboard?

- There is also another function to read from keyboard:

```
char *gets(char *s);
```

- Reads from keyboard until `\n` and stores results in buffer `s` points to
- Why should you never use `gets`?

9

## Example for reading from files

```
#include <stdio.h>
#include <stdlib.h>

#define MAX 100

int main() {
    char buffer[MAX];
    FILE *fp = fopen("my_file.txt", "r");

    if( fp == NULL )
    {
        perror("Error while opening the file.\n");
        exit(1);
    }

    while ((fgets(buffer, MAX, fp)) != NULL) {
        printf("%s", buffer);
    }

    fclose(fp);
}
```

10

## Block I/O (aka binary I/O)

- Suppose a program wants to store a large set of numbers to file for later re-use (not human consumption).
- Could use `fprintf` for that. Why shouldn't we?

- Suppose you store the number 1,999,999 as follows:

```
fprintf(fp, "1999999");
```

- How many bytes does this take up in the file?

```
00110001 00110001 00110001 00110001 00110001 00110001 00110001
```

- How many bytes should it really take to store this number?

```
00011110 10000100 01111111
```

- Block I/O allows you to read and write binary data, i.e. you read and write byte-for-byte rather than lines of characters.

11

## Binary I/O

- Recall that `fgets` reads **characters**.
- By contrast, `fread` and `fwrite` operate on bytes.

```
size_t fread(void *ptr, size_t size,
             size_t nmemb, FILE *stream);
```

- read `nmemb * size` bytes into memory at `ptr`
- returns number of items read

```
size_t fwrite(const void *ptr, size_t size,
             size_t nmemb, FILE *stream);
```

- write `nmemb * size` bytes from `ptr` to the file pointer `stream`
- returns number of items written

12

Example using  
 size\_t fwrite(const void \*ptr, size\_t size,  
 size\_t nmemb, FILE \*stream);

```
/* write an integer to the file */
int num = 1999999;
n = fwrite(&num, sizeof(num), 1, fp);

/* write a struct to the file */
struct rec {
  char name[20];
  int num;
} r;
r.num = 42;
strncpy(r.name, "koala", 20);
n = fwrite(&r, sizeof(r), 1, fp);
```

13

Example using  
 size\_t fread(void \*ptr, size\_t size,  
 size\_t nmemb, FILE \*stream);

```
/* read an integer from the file */
int num;
n = fread(&num, sizeof(num), 1, fp);

/* read a struct from the file */
struct rec r;
n = fread(&r, sizeof(r), 1, fp);

/* display the contents of the variables */
printf("%d %s %d\n", num, r.name, r.num);
```

14

## The plan for today

### • Processes – Unix System Calls

- fork()
- wait()
- exec()


Haviland Chapters 5

15

## Processes can create other processes


- The fork system call creates a **child** process:  

```
#include <unistd.h>
pid_t fork(void);
```
- The fork system call creates a **duplicate** of the currently running program.
- Both processes run concurrently and independently
- It takes no arguments and returns a process ID.



16

## Example: A shell



- When a command is typed, shell forks and then execs the typed command.

17

## What is different between parent and forked child

- The fork system call creates a child process:  

```
#include <unistd.h>
pid_t fork(void);
```
- The child gets a new PID (process ID) and PPID
- The return value from the fork call is different:
  - On success
    - fork() returns 0 to the child
    - fork() returns the child's PID to the parent.
  - On failure (no child created) fork returns -1 to parent

18

## Fork example

```
int main ()
{
    pid_t pid;
    pid = fork();
    if (pid < 0) {
        perror("fork()");
    } else if (pid > 0) {
        printf("parent\n");
    } else { /* pid == 0 */
        printf("child\n");
    }
    return 0;
}
```

19

## Another fork example

Original process (parent)

```
int i; pid_t pid;
i = 5;
printf("%d\n", i);
/* prints 5 */
pid = fork();
/* pid == 677 */
if (pid > 0)
    i = 6;
else (pid == 0)
    i = 4;
printf("%d\n", i);
```

Child process

```
int i; pid_t pid;
i = 5;
printf("%d\n", i);

pid = fork();
/* pid == 0 */
if (pid > 0)
    i = 6;
else if (pid == 0)
    i = 4;
printf("%d\n", i);
```

20


## Yet another fork example

```
int main (void) {
    int i;
    pid_t child;

    for (i = 0; i < 3; i++) {
        child = fork();
        if (child == 0)
            printf ("Child %d saying hi\n", getpid());
    }
    return 0;
}
```


21

## Yet another fork example




22

## Yet another fork example

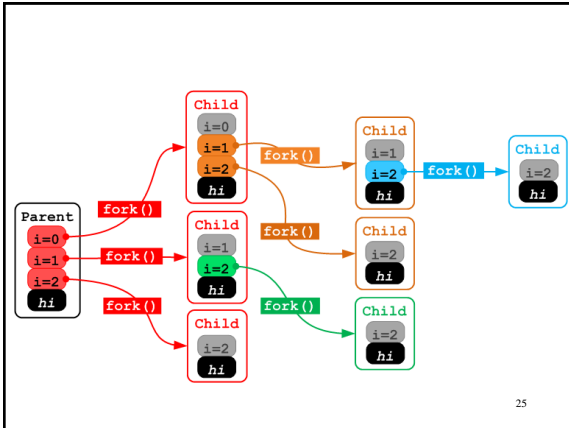


23

## Yet another fork example



24



25

## When does a child terminate?

- Like any other program:
  - The program's main function returns
  - The program calls `exit`
    - `void exit(int status);`
  - The program receives a signal that causes it to terminate
    - Will see next lecture what that means ...
- Remember that programs have an exit status
  - Exit status of most recent command stored in `$_` in most shells
- Exit status is set by call to `exit` or main's `return`

26

## wait()

- A parent might want to wait for a child to complete
- A parent might want to know the exit code of a child
- System call to wait for a child
  - `pid_t wait(int *status)`
- `wait` suspends execution of the calling process until one of its children terminates

27

## wait()

- System call to wait for a child
  - `pid_t wait(int *status)`
- After calling `wait()` a process will:
  - block if all of its children are still running
  - return immediately with the PID of a terminated child, if there is a terminated child
  - return immediately with an error (-1) if it doesn't have any child processes.

28

## Info returned by wait()

- System call to wait for a child
  - `pid_t wait(int *status)`
- Returns the pid of the terminated child or -1 on error
- `status` encodes the exit status of the child and how a child exited (normally or killed by signal)
- There are macros to process exit status:
  - `WIFEXITED` tells you if child terminated normally
  - `WEXITSTATUS` gives you the exit status

29

## Example using wait()

```

int main (void) {
    pid_t child;
    int status, exit_status;

    if ((child = fork()) == 0) {
        sleep (5);
        exit (8);
    }

    wait (&status);
    if (WIFEXITED(status)) {
        exit_status = WEXITSTATUS(status);
        printf ("Child %d done: %d\n", child, exit_status);
    }

    return 0;
}

```

30

## How does a child become a zombie?

- When a child terminates, but its parent process is not waiting for it
- The child (its exit code) is kept around as a zombie until parent collects its exit code through `wait`
  - or until parent terminates
- Shows up as `Z` in `ps`



## How does a child become an orphan?

- If the parent process terminates before the child
- Who is now the new parent?
  - Orphans get adopted by the `init` process
  - `init` is the first process started during booting
  - It's the root of the process hierarchy
  - `init` has a PID of 1
  - The PPID of orphans is 1



## waitpid

- What if a process wants to wait for a particular child (rather than any child)
- What if a process does not want to block when no child has terminated?

```
pid_t waitpid(pid_t pid, int *status, int options);
```

- First parameter specifies PID of child to wait for
- If options is 0, `waitpid` blocks (just like `wait`)
- If options is `WNOHANG`, it immediately returns 0 instead of blocking when no terminated child

33

## Waitpid example

```
int main (void) {
    pid_t child;
    int status, exit_status;
    if ((child = fork()) == 0) {
        sleep (5);
        exit (8);
    }
    while (waitpid (child, &status, WNOHANG) == 0) {
        printf ("Waiting...\n");
        sleep (1);
    }
    if (WIFEXITED(status)) {
        exit_status = WEXITSTATUS(status);
        printf ("Child %d done: %d\n", child, exit_status);
    }
    ...
}
```

34

## Wait a second ....


- Fork creates a duplicate of the current process
- How do we actually create a new process that runs a different program???

*When a program wants to have another program running in parallel, it will typically first use **fork**, then the child process will use **exec** to actually run the desired program.*

35

## Exec

- The `exec` system call replaces the program being run by a process by a different one.
- The new program starts executing from the beginning.
- On success, `exec` never returns, on failure, `exec` returns -1.



36

## Exec example

### Program X

```
int i = 5;
printf("%d\n", i);
exec("Y");
printf("%d\n", i);
```

### Program Y

```
printf("hello\n");
```

37

## exec properties

- New process inherits from calling process:
  - PID and PPID, UID, GID
  - controlling terminal
  - CWD, resource limits
  - pending signals

38

## exec()

- exec() is not one specific function, but a family of functions:

```
execl(char *path, char *arg0, ..., (char *)NULL);
execv(char *path, char *argv[]);
execlp(char *file, char *arg0, ..., (char *)NULL);
execvp(char *file, char *argv[]);
```

- First parameter: name of executable; then commandline parameters for executable; these are passed as argv[0], argv[1], ..., to the main program of the executable.
- execl and execv differ from each other only in how the arguments for the new program are passed
- execlp and execvp differ from execl and execv only in that you don't have to specify full path to new program <sup>39</sup>

39

## Example using execl

```
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>

int main (void) {
    printf ("Before exec\n");
    execl ("/bin/ls", "ls", "-l", (char *)NULL);
    perror ("execl");
    exit (1);
}
```

40


## Example using execv

```
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>

int main (void) {
    char *args[] = {"ls", "-l", NULL};
    printf ("Before exec\n");
    execv ("/bin/ls", args);
    perror ("execv");
    exit (1);
}
```

41

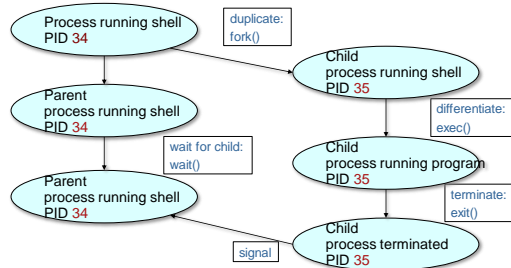
## How a shell runs commands



- When a command is typed, shell forks and then execs the typed command.

42

## How a shell runs commands



43

## Shell skeleton

```

while (1) { // infinite while loop
    print_prompt();
    read_command(command, parameters);
    if (fork()) {
        wait(&status);
    } else {
        execve(command, parameters, NULL);
    }
}

```

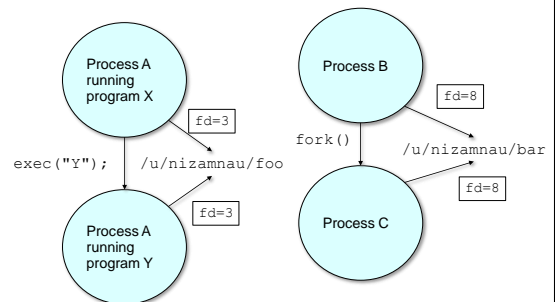
44

## Processes and File Descriptors

- File descriptors are handles to open files.
- They belong to processes not programs.
- They are a process's link to the outside world.


45

## FDs preserved across fork and exec



46

## Initializing Unix



- See "top", "ps -aux" to see what's running
- The **only** way to create a new process is to duplicate an existing process. Therefore the ancestor of **all** processes is `init` with `pid = 1`
- The **only** way to run a program is with `exec`

47

## Shell Job control

- A job (or process) is program in execution
  - Use `ps` to view processes
- Foreground job: has control of the terminal
- Background job: runs concurrently with shell in the background
  - To run a program in the background append `&` to the name of the program
- At any point a program can be running or suspended
  - Hit `<ctrl> z` to suspend the current foreground job

48



## Shell Job control

- `jobs` gives you a list of jobs; each is associated with a job number
- `fg [num]` puts job `num` in the foreground
- `bg [num]` puts job `num` in the background
- `kill %num` kills job `num`
  - You can kill the current foreground job with Ctrl-c
  - You can suspend (pause) current foreground job with Ctrl-z

49

That's it for today!

50